

Chapter 3-4 (chapter 2-3 in online book)

Innleiðing í tilgjörðum viti

Notes

Contents

1	Monte Carlo Tree Search & AlphaGo	3
1.1	AlphaGo Overview	3
1.2	The Game of Go	3
2	Limitations of Traditional Approaches	3
2.1	Minimax with Alpha-Beta Pruning	3
3	Monte Carlo Tree Search (MCTS)	3
3.1	Overview	3
3.2	Core Components	4
3.3	Benefits and Drawbacks	4
4	Exploration vs Exploitation (Multi-Armed Bandit Problem)	4
4.1	Analogy	4
4.2	Upper Confidence Bound (UCB)	4
5	Deep Learning in AlphaGo	5
5.1	Policy and Value Networks	5
5.2	Self-Play and Reinforcement Learning	5
6	Symbolic vs Sub-Symbolic Strategies	5
6.1	Symbolic AI	5
6.2	Sub-Symbolic AI	5
7	Knowledge-Based Agents	5
7.1	Definition	5
7.2	Knowledge Base (KB)	5
7.3	Declarative vs Procedural Approaches	6
8	Logic and Inference	6
8.1	Syntax and Semantics	6
8.2	Logical Inference	6
8.3	Properties	6

9	Propositional Logic	6
9.1	Syntax	6
9.2	Semantics (Models)	6
9.3	Entailment in Propositional Logic	7
10	Inference Algorithms	7
11	Logic and Inference: Advanced Concepts	7
11.1	Complexity of Model Checking	7
11.2	Deductive Systems	7
11.3	Inference Rules	7
11.4	Search Problem Analogy for Theorem Proving	8
11.5	Resolution Method	8
11.6	Proof by Contradiction	8
11.7	Horn Clauses and Definite Clauses	9
11.8	Forward and Backward Chaining	9
11.9	Limitations of Propositional Logic	9
11.10	First-Order Logic (FOL)	9
12	Advanced First-Order Logic and Optimization Techniques	10
12.1	Universal and Existential Quantification	10
12.2	First-Order Logic (FOL) Formulas	10
12.3	Key FOL Techniques	10
12.4	Soundness and Completeness in FOL	10
12.5	Limitations of First-Order Logic	10
12.6	Tautology	11
13	Optimization Concepts	11
13.1	Definition and Objective	11
13.2	Local Search Algorithms	11
13.3	Hill Climbing Variants	11
13.4	Iterative Improvement Techniques	12
13.5	Genetic Algorithms	12
13.6	Key Considerations	12

1 Monte Carlo Tree Search & AlphaGo

1.1 AlphaGo Overview

AlphaGo is an AI developed to play the board game Go at a professional level. It defeated top human players using a combination of techniques:

- Deep Learning
- Monte Carlo Tree Search (MCTS)
- High-Performance Computing

1.2 The Game of Go

- Played on a 19×19 grid \Rightarrow very high branching factor.
- Branching factor: ~ 250 moves on average per turn.
- For comparison, chess has a branching factor of ~ 35 .
- Goal: Surround more territory than the opponent.
- Complexity:
 - Estimated state space: $\sim 10^{100}$
 - Very subtle mechanics \Rightarrow traditional evaluation functions (like in chess) are ineffective.

2 Limitations of Traditional Approaches

2.1 Minimax with Alpha-Beta Pruning

- Works well for chess:
 - Modest branching factor
 - Good heuristic evaluation functions
- In Go:
 - Branching factor is too large
 - Evaluation functions are too subtle to model effectively

3 Monte Carlo Tree Search (MCTS)

3.1 Overview

- Heuristic search algorithm applied to decision trees.
- Common in Go (last ~ 10 years).
- Used when no good evaluation function exists.

3.2 Core Components

- **Selection:** Choose promising moves recursively until a leaf node is reached.
- **Expansion:** Add child nodes (possible future states) at the leaf.
- **Simulation:** Play out the game randomly to the end; record win/loss.
- **Backpropagation:** Propagate results back up the tree; update selection policy.

3.3 Benefits and Drawbacks

Pros:

- Does not require a predefined evaluation function.
- Effective in large, complex spaces.

Cons:

- Requires many simulations.
- Slow for very large games like Go.

4 Exploration vs Exploitation (Multi-Armed Bandit Problem)

4.1 Analogy

- Faced with multiple slot machines ("arms").
- Each has an unknown reward distribution.
- Must balance:
 - **Exploitation:** Use the arm that seems best.
 - **Exploration:** Try others to discover potential improvements.

4.2 Upper Confidence Bound (UCB)

- Balances exploration and exploitation.
- Formula considers both:
 - Average reward of each arm.
 - Number of times each arm has been tried.
- Policy:
 - Try each option once.
 - Then choose the option maximizing the UCB formula.

5 Deep Learning in AlphaGo

5.1 Policy and Value Networks

- **Policy Network:** Suggests promising moves, reducing branching factor.
- **Value Network:** Estimates game state value without full simulation.

5.2 Self-Play and Reinforcement Learning

- AlphaGo trained by playing against itself.
- Reinforcement learning improved decision-making, reducing search depth and branching factor.

6 Symbolic vs Sub-Symbolic Strategies

6.1 Symbolic AI

- Based on explicit knowledge and logical reasoning.
- Uses rules, symbols, and logic.
- Examples: Knowledge bases, expert systems, logic inference.

6.2 Sub-Symbolic AI

- Uses numerical/statistical methods.
- Examples: Neural networks, genetic algorithms, deep learning.
- Learns patterns from data rather than explicit rules.

7 Knowledge-Based Agents

7.1 Definition

Agents using internal knowledge representation to reason and derive conclusions.

7.2 Knowledge Base (KB)

- Repository of information stored in sentences.
- Sentences: Assertions about the world in a formal language.

7.3 Declarative vs Procedural Approaches

- **Declarative:** Feed system facts; it reasons to conclusions.
- **Procedural:** Encode specific behaviours directly.
- Often combined in modern systems.

8 Logic and Inference

8.1 Syntax and Semantics

- **Syntax:** Rules for forming valid sentences.
- **Semantics:** Meaning of sentences; correspondence with the world.

8.2 Logical Inference

- Derive new sentences from known ones using rules.
- Goal: Use facts to answer queries or deduce new knowledge.

8.3 Properties

- **Entailment:** α is entailed by KB if α is true in all models where KB is true. Example: $\text{KB} = \{\text{"All humans are mortal"}, \text{"Socrates is a human"}\} \Rightarrow \text{"Socrates is mortal"}$
- **Completeness:** An inference algorithm is complete if it can derive all entailed sentences.

9 Propositional Logic

9.1 Syntax

- Atomic symbols: P, Q, \dots
- Compound sentences:
 - Conjunction: $P \wedge Q$
 - Disjunction: $P \vee Q$
 - Implication: $P \rightarrow Q$

9.2 Semantics (Models)

- A model assigns truth values to atomic propositions.
- A model = one possible world configuration.

9.3 Entailment in Propositional Logic

- $KB \models \alpha \iff \alpha$ is true in all models where KB is true.
- **Truth Table:** Lists all possible truth assignments; used to check entailment.

10 Inference Algorithms

- Goal: Decide if $KB \models \alpha$
- Methods:
 - Truth table checking (exhaustive).
 - More efficient algorithms (e.g., resolution).

11 Logic and Inference: Advanced Concepts

11.1 Complexity of Model Checking

- Determining whether $KB \models Q$ (entailment) can be computationally expensive.
- Naive approach: test all possible interpretations \rightarrow exponential in number of propositions.
- Efficient methods rely on syntactic manipulation rather than enumerating all models.

11.2 Deductive Systems

- **Soundness:** If $KB \vdash Q$ then $KB \models Q$ (only derives truths).
- **Completeness:** If $KB \models Q$ then $KB \vdash Q$ (can derive everything entailed).
- Deduction theorem connects semantic entailment with syntactic derivation.

11.3 Inference Rules

- **Modus Ponens:** If $\alpha \rightarrow \beta$ and α are true, then β is true.
- **Conjunction Elimination:** If $\alpha \wedge \beta$ is true, then α is true.
- **Double Negation Elimination:** $\neg\neg\alpha \equiv \alpha$.
- **Implication Elimination:** $\alpha \rightarrow \beta \equiv \neg\alpha \vee \beta$.

- **Biconditional Elimination:** $\alpha \leftrightarrow \beta \equiv (\alpha \rightarrow \beta) \wedge (\beta \rightarrow \alpha)$.
- **De Morgan's Laws:** $\neg(\alpha \wedge \beta) \equiv \neg\alpha \vee \neg\beta$, $\neg(\alpha \vee \beta) \equiv \neg\alpha \wedge \neg\beta$.
- **Distributive Laws:** $\alpha \wedge (\beta \vee \gamma) \equiv (\alpha \wedge \beta) \vee (\alpha \wedge \gamma)$, etc.

11.4 Search Problem Analogy for Theorem Proving

- **Initial State:** Knowledge Base (KB)
- **Actions:** Inference rules applied to current KB
- **Transition Model:** New KB after inference
- **Goal Test:** Statement Q to prove
- **Path Cost:** Number of steps or inferences used

11.5 Resolution Method

- **Principle:** Combine clauses to derive new information until goal or contradiction.
- Example of resolving literals:

$$P \vee Q, \quad \neg P \Rightarrow Q$$

$$P \vee Q, \quad \neg P \vee R \Rightarrow Q \vee R$$

- **Clause:** A disjunction of literals, e.g., $P \vee Q \vee R$.
- **Conjunctive Normal Form (CNF):** Conjunction of clauses, e.g., $(A \vee B \vee C) \wedge (D \vee \neg E)$.
- Conversion to CNF:
 - Eliminate biconditionals and implications
 - Move NOT inwards using De Morgan's laws
 - Apply distributive laws

11.6 Proof by Contradiction

- Assume $\neg Q$ and derive a contradiction from KB.
- Ensures KB consistency.

11.7 Horn Clauses and Definite Clauses

- Horn Clause: CNF with at most one positive literal. Form: $\neg P \vee Q \equiv P \rightarrow Q$
- Definite Clause: Exactly one positive literal (the "head"), possibly many negative literals (the "body").
- Basis for efficient logic programming and automated reasoning.

11.8 Forward and Backward Chaining

- **Forward Chaining:** Start from known facts, apply inference rules to derive new facts until goal.
- **Backward Chaining:** Start from query Q , work backwards applying rules until known facts are reached.
- Linear resolution strategies (e.g., SLD resolution) reduce search space using sub-goals.

11.9 Limitations of Propositional Logic

- Cannot represent objects, relations, or quantifiers.
- Limited expressiveness for complex domains.

11.10 First-Order Logic (FOL)

- Extends propositional logic with:
 - Constants: specific objects
 - Variables: placeholders
 - Predicates: properties/relations over objects
 - Functions: map objects to objects
- Terms: constants, variables, or functions applied to terms
- Atomic formulas: predicates applied to terms
- Quantifiers:
 - Universal: $\forall x P(x)$
 - Existential: $\exists x P(x)$

12 Advanced First-Order Logic and Optimization Techniques

12.1 Universal and Existential Quantification

- **Universal Quantification:** $\forall x P(x)$ — "for all x , $P(x)$ is true".
- **Existential Quantification:** $\exists x P(x)$ — "there exists an x such that $P(x)$ is true".
- These quantifiers are fundamental in First-Order Logic (FOL) for expressing general statements about objects.

12.2 First-Order Logic (FOL) Formulas

- **Arity:** Number of arguments a predicate or function takes.
- **Predicates with functions:** Allow expressing relationships between objects and their properties.
- **Interpretations:** Map constants, functions, and predicates to values or truth assignments. Example: $F = c_1 + c_3 > c_2$, domain $\{1, 2, 3, 4\}$ Relations are sets of tuples where the interpretation is true.

12.3 Key FOL Techniques

- **Skolemization:** Remove existential quantifiers by introducing Skolem constants or functions.
- **Unification:** Process of finding substitutions for variables to make terms or predicates match.
- **Resolution in FOL:** Extends propositional resolution to FOL using unification to derive new clauses.

12.4 Soundness and Completeness in FOL

- **Soundness:** Every derived formula is logically true.
- **Completeness:** Every logically entailed formula can be derived using the inference system.

12.5 Limitations of First-Order Logic

- Incompleteness in some domains
- Scalability challenges with large knowledge bases
- Difficult to represent spatial and temporal reasoning
- Handling uncertainty is non-trivial

12.6 Tautology

- A formula that is true under every possible interpretation.
- Important in both propositional logic and FOL.

13 Optimization Concepts

13.1 Definition and Objective

- Optimization: Selecting the best option from a set of alternatives.
- Formalization:
 - Variables with associated domains
 - Objective function mapping assignments to real numbers
 - Optimality criterion: find assignment minimizing or maximizing the objective function
- Example: Minimizing loss associated with a variable assignment.

13.2 Local Search Algorithms

- Maintain a single current state and explore neighboring states to improve the objective.
- **Example:** Placing houses and hospitals in 2D space to minimize total distance.
- **State space landscape:** Visualizes global/local maxima and minima, flat regions, and shoulders.

13.3 Hill Climbing Variants

- **Steepest-Ascent Hill Climbing:** Move to the neighbor with the highest improvement.
 - Issues: Local maxima, plateaus, ridges
- **Stochastic Hill Climbing:** Randomly select among improving neighbors.
- **First-Choice Hill Climbing:** Randomly evaluate neighbors and move to first improvement found.
- **Random-Restart Hill Climbing:** Restart from random initial states to escape local maxima.
- **Local Beam Search:** Maintain k states in parallel, keep best successors at each step.

13.4 Iterative Improvement Techniques

- **Simulated Annealing:** Accept worse neighbors with decreasing probability over time to escape local optima.
- **Gradient Descent:** For continuous domains, adjust variables proportionally to reduce the objective function.

13.5 Genetic Algorithms

- Maintain a population of candidate solutions.
- Randomly select pairs and perform crossover to produce offspring.
- Apply mutation to introduce variability.
- Iterate until an acceptable solution is found.

13.6 Key Considerations

- Local search is fast but can get stuck in local optima.
- Random restarts, stochastic methods, and population-based methods improve robustness.
- Continuous vs discrete domains may require specialized algorithms (e.g., gradient descent for continuous).